

ATISHAY JAIN

FULL-STACK SOFTWARE ENGINEER

Ludhiana, Punjab, India | +91 9646009828 | ajain8_be23@thapar.edu | LinkedIn | GitHub | Portfolio

EDUCATION

B.Tech, Computer Engineering - Thapar Institute of Engineering and Technology | Aug 2023 - Jun 2027

CGPA: 8.0/10.0 | Patiala, Punjab, India

EXPERIENCE & LEADERSHIP

Def-Space Summer Internship Programme - Bharat Space Education Research Centre (BSERC) | Jun 2026 - Aug 2026

- Completed a 120-hour, 6+ week Def-Space programme focused on defence and space technologies and responsible innovation.

Finance & Documentation Manager - TEDxTIET | Aug 2025 - Feb 2026

- Managed a INR 3.5 lakh event budget through allocation, expense tracking, sponsor records, and variance reporting.

PROJECTS

ReceiptSplit - Multi-room Bill-Splitting Platform | 2026

Next.js, TypeScript, FastAPI, async SQLAlchemy, PostgreSQL, Docker, SSE, OCR

- Built a multi-room bill-splitting MVP with OCR drafts, deterministic split rounding, durable room events, and stateful settlement workflows.
- Expanded frontend regression checks from 51 to 82 (+61%) and maintained 379 backend test functions across 53 files for authenticated room/member and settlement flows.

TIET-LOOP - Campus Event Platform | 2025 - 2026

React, Vite, Tailwind CSS, FastAPI, WebSockets, Google Maps

- Collaborative Full-Stack Contributor to a campus platform with 8 persistence models spanning authentication, profiles, discovery, RSVPs, friends, chat, carpool, maps, recommendations, and admin workflows.
- Contributed React/FastAPI integration and real-time WebSocket flows; signed JWT access tokens expire after 30 minutes and passwords use Argon2 hashing.

Hybrid GenAI Transaction Categorizer | 2025 - 2026

Python, FastAPI, React, ONNX Runtime, DistilBERT, SQLite, local LLM integration

- Built a 3-path offline categorization pipeline - merchant rules, ONNX DistilBERT, then local-LLM fallback - over 1,600 labeled transaction rows across 10 categories.
- Reduced local CPU fallback latency from 60-80 seconds to 5-13 seconds (78-94% lower) by consolidating Qwen2.5, shortening prompts, and capping output at 40 tokens; passed 27 local tests.

Mahoraga - Adaptive RL Environment | 2026

Python, Gymnasium, FastAPI, Gradio, Qwen 2.5, LoRA

- Contributed as a Backend / RL Systems Engineer on a 5-action adaptive environment with inspectable reward components, Gymnasium support, and LLM/backend integration.
- Meta OpenEnv Hackathon 2026 Finalist; sped up T4 fine-tuning by ~3x (1,024 to 512 token context; batch size 2 to 4) and verified 145 Python checks.

SKILLS

Languages: Python, TypeScript, JavaScript, C++, C, SQL

Full-Stack: React, Next.js, Vite, Tailwind CSS, FastAPI, REST APIs, WebSockets, JWT, Google Maps

Data & Infrastructure: PostgreSQL, Supabase, SQLite, SQLAlchemy, Docker, Git, GitHub Actions, Render

Applied AI: ONNX Runtime, DistilBERT, Transformers, local LLM integration, Gymnasium, LoRA

CERTIFICATIONS

BSERC Def-Space Summer Internship Programme (120 hours) | NVIDIA Foundations of Deep Learning | Duke University -

Introduction to RAG | CI/CD Pipeline with Docker